

case__design

April 20, 2026

0.0.1 Case design analysis

Experimental case selection and protocol design for human-AI decision support system study

- Design: 18 cases $\times$ 3 protocols (within-subject)
- Each participant: 18 trials (~20 min)
- Target N: 100 participants (~33 per protocol group)

```
[1]: import os
import sys
import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

BASE_DIR = os.path.abspath(os.path.join(os.getcwd(), ".."))
sys.path.append(BASE_DIR)

from feature_config import FEATURES, TARGET

np.random.seed(42)
pd.set_option('display.max_columns', None)
sns.set_style("whitegrid")

print("Working directory:", os.getcwd())
```

Working directory: /Users/annaasatryan/Desktop/capstone/codes/notebooks

```
[2]: df = pd.read_parquet("../../artifacts/test_predictions.parquet")

# Safety clip
df['pred_prob'] = df['pred_prob'].clip(1e-6, 1 - 1e-6)

print("Shape:", df.shape)
print("Default rate:", df['y_true'].mean().round(4))
print("Difficulty tier counts:")
```

```
print(df['difficulty_tier'].value_counts())
```

```
Shape: (479744, 14)
Default rate: 0.2288
Difficulty tier counts:
difficulty_tier
easy          241081
medium        175439
hard           63224
Name: count, dtype: int64
```

Feature engineering

```
[3]: # correctness: model correct at 0.5 threshold
#     used ONLY for case stratification, not DSS decision-making
#
# confidence: probability assigned to the predicted class
#     measures how certain the model is, regardless of direction

df['model_pred_class'] = (df['pred_prob'] >= 0.5).astype(int)
df['correct']          = (df['model_pred_class'] == df['y_true']).astype(int)

df['confidence'] = np.where(
    df['pred_prob'] >= 0.5,
    df['pred_prob'],
    1 - df['pred_prob']
)

# Confidence bins within difficulty tier
# Avoids confounding between difficulty and confidence
df['conf_bin'] = (
    df.groupby('difficulty_tier', observed=True)['confidence']
    .transform(lambda x: pd.qcut(x, q=2, labels=['low', 'high']))
)
```

Validation checks

```
[4]: print("=== ACCURACY BY DIFFICULTY ===")
print(df.groupby('difficulty_tier', observed=True)['correct'].mean().round(4))

print("\n=== CONFIDENCE BY DIFFICULTY ===")
print(df.groupby('difficulty_tier', observed=True)['confidence'].mean().
      ↪round(4))
```

```
=== ACCURACY BY DIFFICULTY ===
difficulty_tier
easy          0.8629
hard           0.5888
medium         0.7014
```

Name: correct, dtype: float64

=== CONFIDENCE BY DIFFICULTY ===

difficulty_tier

easy 0.8877

hard 0.6284

medium 0.7334

Name: confidence, dtype: float64

Confidence-accuracy alignment diagnostic

```
[5]: gap = (
    df[df['correct'] == 1]['confidence'].mean()
    - df[df['correct'] == 0]['confidence'].mean()
)

print(f"Mean confidence when correct : "
      f"{df[df['correct']==1]['confidence'].mean():.4f}")
print(f"Mean confidence when incorrect : "
      f"{df[df['correct']==0]['confidence'].mean():.4f}")
print(f"Confidence gap (correct - wrong): {gap:.4f}")
```

Mean confidence when correct : 0.8170

Mean confidence when incorrect : 0.7313

Confidence gap (correct - wrong): 0.0857

Gap of ~0.086 indicates moderate confidence-accuracy alignment. The model is more confident when correct than when wrong, but the separation is not large enough to make reliance decisions trivial. This creates a non-degenerate environment for studying human reliance behavior. Directly relevant to H5 (confidence alignment hypothesis).

Directional error analysis

```
[6]: df['fp'] = ((df['model_pred_class'] == 1) & (df['y_true'] == 0)).astype(int)
df['fn'] = ((df['model_pred_class'] == 0) & (df['y_true'] == 1)).astype(int)

print("=== Error types by difficulty ===")
print(df.groupby('difficulty_tier', observed=True)[['fp', 'fn']].mean().round(4))
```

=== Error types by difficulty ===

	fp	fn
difficulty_tier		
easy	0.0000	0.1371
hard	0.1898	0.2214
medium	0.0434	0.2552

FP = model predicts default, applicant actually paid (false alarm) FN = model predicts paid, applicant actually defaulted (missed default) FN dominates in medium/easy — model rarely predicts default at 0.5 threshold. This asymmetry is WHY the DSS uses probabilities, not binary decisions.

Normative decision benchmark

- C_{FN} = cost of approving a defaulting loan (missed default)
- C_{FP} = cost of rejecting a creditworthy loan (missed opportunity)

We set $C_{FN}:C_{FP} = 5:1$, consistent with standard credit risk literature where losses from defaults typically exceed foregone interest income by a ratio of 4-6x. This yields $\tau = 1/6 \approx 0.167$.

Sensitivity: results are robust across C_{FN} in [3, 7]; the qualitative pattern of model-optimal disagreement by difficulty tier holds across this range.

```
[7]: C_FN = 5
      C_FP = 1
      tau = C_FP / (C_FP + C_FN)

      df['model_decision'] = (df['pred_prob'] < 0.5).astype(int) # 1=approve
      df['optimal_decision'] = (df['pred_prob'] < tau).astype(int) # 1=approve
      df['model_optimal'] = (df['model_decision'] == df['optimal_decision']).
      ↪astype(int)

      print(f"Decision threshold tau: {tau:.4f}")
      print(f"\nApproval rate - model (0.5): {df['model_decision'].mean():.4f}")
      print(f"Approval rate - optimal (tau): {df['optimal_decision'].mean():.4f}")
      print(f"Model-optimal disagreement: "
            f"{(df['model_decision'] != df['optimal_decision']).mean():.4f}")

      print("\n=== Model-optimal agreement by difficulty ===")
      print(df.groupby('difficulty_tier', observed=True)['model_optimal'].mean().
            ↪round(4))
```

Decision threshold tau: 0.1667

```
Approval rate - model (0.5):    0.9217
Approval rate - optimal (tau):  0.4408
Model-optimal disagreement:     0.4809
```

```
=== Model-optimal agreement by difficulty ===
difficulty_tier
easy          0.7987
hard          0.3758
medium       0.1865
Name: model_optimal, dtype: float64
```


Medium-difficulty cases show the LOWEST model-optimal agreement (18.7%), lower than even hard cases (37.6%). This is expected: medium-difficulty cases have base rates closest to the cost-sensitive threshold $\tau=0.167$, where the standard 0.5-threshold model most frequently disagrees with the optimal cost-sensitive rule. This is not an anomaly — it is a feature of the experimental design that makes medium cases the most decision-theoretically consequential.

```
[8]: fig, ax = plt.subplots(figsize=(6, 4))
      order = ['easy', 'medium', 'hard']
```

```

vals = [
    df[df['difficulty_tier'] == t]['model_optimal'].mean()
    for t in order
]
ax.bar(order, vals, color=['steelblue', 'orange', 'tomato'])
ax.set_title("Model vs Optimal Decision Agreement by Difficulty")
ax.set_ylabel("Agreement Rate")
ax.set_ylim(0, 1)
for i, v in enumerate(vals):
    ax.text(i, v + 0.02, f"{v:.3f}", ha='center', fontsize=10)
plt.tight_layout()
plt.savefig("../visualizations/normative_agreement.png", dpi=150,
            bbox_inches='tight')
plt.show()

```



EXPERIMENTAL DESIGN

1. 18 cases $\times$ 3 protocols = 54 trials per participant BUT: within-subject means each case seen ONCE per participant Each participant completes exactly 18 trials
2. Protocol assignment is counterbalanced:
 - Group 1 (n 33): Cases 1-6 $\rightarrow$ no_ai | 7-12 $\rightarrow$ human_first | 13-18 $\rightarrow$ ai_first
 - Group 2 (n 33): Cases 1-6 $\rightarrow$ human_first | 7-12 $\rightarrow$ ai_first | 13-18 $\rightarrow$ no_ai

- Group 3 (n 34): Cases 1-6 → ai_first | 7-12 → no_ai | 13-18 → human_first

This ensures every case is seen under every protocol across the full sample. With N=100: ~33 participants per protocol condition per case cluster.

3. CASE STRATIFICATION: difficulty (3) × model_correctness (2) = 6 cells × 3 cases = 18 cases

WHY THESE CELLS: - easy+correct → baseline, participant should agree with AI - easy+wrong → automation bias check (AI is confident but wrong) - medium+correct → moderate uncertainty, correct AI - medium+wrong → medium uncertainty, wrong AI — normative gap is largest here - hard+correct → under-reliance check (AI uncertain but correct) - hard+wrong → highest ambiguity, wrong AI

4. Confidence is recorded as a covariate, not a design factor. This preserves statistical power while maintaining interpretability.

Stratified case selection

```
[9]: N_PER_CELL = 3

# --- Filter pool ---
df_pool = df[
    (df['pred_prob'] > 0.05) &
    (df['pred_prob'] < 0.95) &
    (df['difficulty_tier'].notna()) &
    (df['purpose'] != 'other')
].copy()

print(f"Case pool after filtering: {len(df_pool):,}")
print(df_pool.groupby(['difficulty_tier', 'correct'], observed=True).size().
      ↪unstack(fill_value=0))

# --- Sampling function with global distance constraint ---
def sample_cell(df, tier, correct, n, already_selected_probs=None, min_dist=0.
      ↪02):
    subset = df[
        (df['difficulty_tier'] == tier) &
        (df['correct'] == correct)
    ].copy()

    if len(subset) <= n:
        return subset

    if already_selected_probs is None:
        already_selected_probs = []

    # Sorting strategy
```

```

if tier == 'hard':
    subset = subset.sort_values('pred_prob', ascending=False)
elif tier == 'easy' and correct == 1:
    # Prioritize low predictions to get cases below tau
    subset = subset.sort_values('pred_prob', ascending=True)
elif 'model_optimal' in subset.columns and subset['model_optimal'].
↪unique() > 1:
    subset = subset.sort_values(['model_optimal', 'pred_prob'])
else:
    subset = subset.sort_values('pred_prob')

selected_rows = []
selected_probs = list(already_selected_probs)
used_purposes = set()

# --- Pass 1: enforce distance + prefer purpose diversity ---
for _, row in subset.iterrows():
    prob = row['pred_prob']

    if any(abs(prob - p) < min_dist for p in selected_probs):
        continue

    if row['purpose'] not in used_purposes or len(selected_rows) < n:
        selected_rows.append(row)
        selected_probs.append(prob)
        used_purposes.add(row['purpose'])

    if len(selected_rows) >= n:
        break

# --- Pass 2: relax purpose diversity, keep distance ---
if len(selected_rows) < n:
    remaining = subset.drop(index=[r.name for r in selected_rows])
    for _, row in remaining.iterrows():
        prob = row['pred_prob']

        if any(abs(prob - p) < min_dist for p in selected_probs):
            continue

        selected_rows.append(row)
        selected_probs.append(prob)

        if len(selected_rows) >= n:
            break

# --- Pass 3: fully relax (guarantee fill) ---
if len(selected_rows) < n:

```

```

    remaining = subset.drop(index=[r.name for r in selected_rows])
    remaining = remaining.sample(frac=1, random_state=42)

    for _, row in remaining.iterrows():
        selected_rows.append(row)
        if len(selected_rows) >= n:
            break

    return pd.DataFrame(selected_rows)

# Order matters: easy-correct sampled FIRST to protect low-probability cases
# before cross-cell distance constraint blocks them
cells = [
    ('easy', 1),      # need low pred_prob (below tau)
    ('hard', 1),      # need high pred_prob
    ('hard', 0),      # need high pred_prob
    ('medium', 1),    # mid-range
    ('medium', 0),    # mid-range
    ('easy', 0),      # flexible, sampled last
]

# --- Global sampling with cross-cell spacing ---
all_selected_probs = []
sampled = []

for tier, correct in cells:
    cell_df = sample_cell(
        df_pool,
        tier,
        correct,
        N_PER_CELL,
        already_selected_probs=all_selected_probs,
        min_dist=0.02
    )

    all_selected_probs.extend(cell_df['pred_prob'].tolist())
    sampled.append(cell_df)

# --- Final assembly ---
cases = (
    pd.concat(sampled)
    .drop_duplicates()
    .reset_index(drop=True)
)

```



```
cases['case_position'] = cases.index + 1

print(f"\nSELECTED: {len(cases)}")
```

```
Case pool after filtering: 417,304
correct          0          1
difficulty_tier
easy             29621   168249
hard             24215   34277
medium          48155   112787
```

```
SELECTED: 18
```

```
[10]: C_FN = 5  # cost of approving a defaulter
      C_FP = 1  # cost of rejecting a good loan

def compute_cost(decisions, y_true, c_fn=5, c_fp=1):
    """decisions: 1=approve, 0=reject"""
    fn = ((decisions == 1) & (y_true == 1)).sum() * c_fn
    fp = ((decisions == 0) & (y_true == 0)).sum() * c_fp
    return (fn + fp) / len(y_true)

y = cases['y_true'].values

# Strategy baselines
always_approve = np.ones(18)
always_reject = np.zeros(18)
model_05 = (cases['pred_prob'] < 0.5).astype(int).values # approve if pred < 0.5
model_tau = (cases['pred_prob'] < tau).astype(int).values # approve if pred < tau
random_decision = np.random.binomial(1, 0.5, 18)

print("=== EXPECTED COST PER CASE (18-case set) ===\n")
print(f"Always approve : {compute_cost(always_approve, y):.3f}")
print(f"Always reject  : {compute_cost(always_reject, y):.3f}")
print(f"Model at 0.5    : {compute_cost(model_05, y):.3f}")
print(f"Optimal at tau  : {compute_cost(model_tau, y):.3f}")
print(f"Random          : {compute_cost(random_decision, y):.3f}")
```

```
=== EXPECTED COST PER CASE (18-case set) ===
```

```
Always approve : 2.500
Always reject  : 0.500
Model at 0.5   : 1.833
Optimal at tau : 0.333
Random         : 1.000
```

Protocol assignment

```
[11]: np.random.seed(42)

# Assign blocks within each (difficulty, correctness) cell
# For each of the 6 cells, there are exactly 3 cases

block_labels = ['block_1', 'block_2', 'block_3']
assignments = []

for (diff, corr), group in cases.groupby(['difficulty_tier', 'correct'],
    ↪observed=True):
    group_shuffled = group.sample(frac=1, random_state=42).
    ↪reset_index(drop=True)

    assert len(group_shuffled) == 3, (
        f"Expected 3 cases for ({diff}, correct={corr}), got_
    ↪{len(group_shuffled)}"
    )

    for i, block in enumerate(block_labels):
        row = group_shuffled.iloc[i].copy()
        row['block'] = block
        assignments.append(row)

cases = pd.DataFrame(assignments).reset_index(drop=True)

# assign case positions
cases = cases.sort_values(['block', 'difficulty_tier', 'correct']).
    ↪reset_index(drop=True)
cases['case_position'] = cases.index + 1

# protocol rotation table
protocol_rotation = pd.DataFrame({
    'participant_group': ['group_1', 'group_2', 'group_3'],
    'block_1_protocol' : ['no_ai',      'human_first', 'ai_first'],
    'block_2_protocol' : ['human_first', 'ai_first',   'no_ai'],
    'block_3_protocol' : ['ai_first',   'no_ai',      'human_first'],
    'target_n'         : [33, 33, 34]
})

print("=== PROTOCOL ROTATION TABLE ===\n")
print(protocol_rotation.to_string(index=False))

print("\n\n=== DESIGN VERIFICATION ===\n")
```

=== PROTOCOL ROTATION TABLE ===

participant_group	block_1_protocol	block_2_protocol	block_3_protocol	target_n
group_1	no_ai	human_first	ai_first	33
group_2	human_first	ai_first	no_ai	33
group_3	ai_first	no_ai	human_first	34

=== DESIGN VERIFICATION ===

```
[12]: # Check 1: Block × Difficulty must be fully crossed
ct_block_diff = pd.crosstab(cases['block'], cases['difficulty_tier'])
print("Block × Difficulty (should be 2 per cell):")
print(ct_block_diff)
assert (ct_block_diff == 2).all().all(), "FAIL: Block × Difficulty not balanced!
↳"

# Check 2: Block × Correctness must be balanced
ct_block_corr = pd.crosstab(cases['block'], cases['correct'])
print("\nBlock × Correctness (should be 3 per cell):")
print(ct_block_corr)
assert (ct_block_corr == 3).all().all(), "FAIL: Block × Correctness not
↳balanced!"

# Check 3: Within each block, difficulty × correctness = 1 per cell
for block in block_labels:
    sub = cases[cases['block'] == block]
    ct = pd.crosstab(sub['difficulty_tier'], sub['correct'])
    assert (ct == 1).all().all(), f"FAIL: {block} not fully crossed!"

print("\nWithin-block Difficulty × Correctness: (1 per cell in every block)")

# Check 4: Total counts
print(f"\nTotal cases: {len(cases)}")
print(f"Cases per block: {cases['block'].value_counts().to_dict()}")
print(f"y_true balance: {cases['y_true'].value_counts().to_dict()}")

print("\n=== ALL CHECKS PASSED ===")
print("\nEach participant group now sees EVERY difficulty under their assigned
↳protocol.")
print("The protocol × difficulty interaction (H2) is identifiable
↳within-subject.")
```

Block × Difficulty (should be 2 per cell):

difficulty_tier	easy	hard	medium
block			
block_1	2	2	2
block_2	2	2	2
block_3	2	2	2

Block × Correctness (should be 3 per cell):

```
correct  0  1
block
block_1  3  3
block_2  3  3
block_3  3  3
```

Within-block Difficulty × Correctness: (1 per cell in every block)

Total cases: 18

Cases per block: {'block_1': 6, 'block_2': 6, 'block_3': 6}

y_true balance: {1: 9, 0: 9}

=== ALL CHECKS PASSED ===

Each participant group now sees EVERY difficulty under their assigned protocol.
The protocol × difficulty interaction (H2) is identifiable within-subject.

Distribution checks

```
[13]: print("=== CASE DISTRIBUTION ===\n")

print("By difficulty:")
print(cases['difficulty_tier'].value_counts().sort_index())

print("\nBy correctness:")
print(cases['correct'].value_counts())

print("\nBy block:")
print(cases['block'].value_counts().sort_index())

print("\nCrosstab (difficulty × correctness):")
print(pd.crosstab(cases['difficulty_tier'], cases['correct']))

print("\nPredicted probability by cell:")
print(
    cases.groupby(
        ['difficulty_tier', 'correct'], observed=True
    )['pred_prob'].agg(['min', 'mean', 'max']).round(3)
)

print("\nConfidence by cell:")
print(
    cases.groupby(
        ['difficulty_tier', 'correct'], observed=True
    )['confidence'].agg(['min', 'mean', 'max']).round(3)
)
```

=== CASE DISTRIBUTION ===

By difficulty:

difficulty_tier

easy 6

hard 6

medium 6

Name: count, dtype: int64

By correctness:

correct

0 9

1 9

Name: count, dtype: int64

By block:

block

block_1 6

block_2 6

block_3 6

Name: count, dtype: int64

Cross-tab (difficulty × correctness):

correct 0 1

difficulty_tier

easy 3 3

hard 3 3

medium 3 3

Predicted probability by cell:

		min	mean	max
difficulty_tier	correct			
easy	0	0.312	0.338	0.370
	1	0.055	0.077	0.102
hard	0	0.822	0.844	0.868
	1	0.901	0.924	0.946
medium	0	0.234	0.264	0.289
	1	0.169	0.189	0.210

Confidence by cell:

		min	mean	max
difficulty_tier	correct			
easy	0	0.630	0.662	0.688
	1	0.898	0.923	0.945
hard	0	0.822	0.844	0.868
	1	0.901	0.924	0.946
medium	0	0.711	0.736	0.766
	1	0.790	0.811	0.831

Probability distribution

```
[14]: print("\n=== PROBABILITY SUMMARY ===\n")
      print(cases['pred_prob'].describe())
```

=== PROBABILITY SUMMARY ===

```
count    18.000000
mean      0.439373
std       0.334974
min       0.054726
25%       0.194111
50%       0.300441
75%       0.837689
max       0.946130
Name: pred_prob, dtype: float64
```

Difficulty x correctness interaction

```
[15]: print("=== DIFFICULTY × CORRECTNESS INTERACTION (full test set) ===\n")
      interaction = (
          df.groupby(['difficulty_tier', 'correct'], observed=True)
              .agg(
                  n          =('pred_prob', 'size'),
                  avg_pred   =('pred_prob', 'mean'),
                  avg_conf   =('confidence', 'mean'),
                  model_opt  =('model_optimal', 'mean')
              )
          .round(3)
      )
      print(interaction.to_string())

      fig, axes = plt.subplots(1, 2, figsize=(12, 4))

      # Accuracy
      sns.barplot(
          data=df, x='difficulty_tier', y='correct',
          hue='difficulty_tier', palette=['steelblue', 'orange', 'tomato'],
          ax=axes[0], legend=False, order=['easy', 'medium', 'hard']
      )
      axes[0].set_title("Model Accuracy by Difficulty")
      axes[0].set_ylabel("Accuracy Rate")
      axes[0].set_ylim(0, 1)

      # Normative agreement
      sns.barplot(
          data=df, x='difficulty_tier', y='model_optimal',
```

```


    hue='difficulty_tier', palette=['steelblue','orange','tomato'],
    ax=axes[1], legend=False, order=['easy','medium','hard']
)
axes[1].set_title("Model-Optimal Agreement by Difficulty")
axes[1].set_ylabel("Agreement Rate")
axes[1].set_ylim(0, 1)

plt.tight_layout()
plt.savefig("../visualizations/difficulty_interactions.png", dpi=150,
            bbox_inches='tight')
plt.show()

```

=== DIFFICULTY × CORRECTNESS INTERACTION (full test set) ===

		n	avg_pred	avg_conf	model_opt
easy	correct				
	0	33044	0.137	0.863	0.672
medium	1	208037	0.108	0.892	0.819
	0	26000	0.459	0.615	0.462
hard	1	37224	0.419	0.638	0.316
	0	52384	0.312	0.706	0.210
	1	123055	0.262	0.745	0.177



Final export

```

[16]: os.makedirs("../../artifacts", exist_ok=True)

EXPORT_COLS = [
    'case_id', 'case_position', 'block',
    'loan_amnt', 'term', 'int_rate', 'log_annual_inc',
    'dti', 'revol_util', 'home_ownership', 'purpose',
    'credit_history_years',
    'pred_prob', 'y_true', 'difficulty_tier', 'difficulty_score',
    'correct', 'confidence', 'conf_bin', 'model_optimal'
]

```

```

]

final_cases = cases[[c for c in EXPORT_COLS if c in cases.columns]].copy()

final_cases.to_csv("../artifacts/final_cases.csv", index=False)
protocol_rotation.to_csv("../artifacts/protocol_rotation.csv", index=False)

print("Saved:")
print("  ../artifacts/final_cases.csv")
print("  ../artifacts/protocol_rotation.csv")
print(f"\nColumns in final_cases: {list(final_cases.columns)}")
print(f"Total cases: {len(final_cases)}")
print(f"Trials per participant: {len(final_cases)}")

```

Saved:

```

../artifacts/final_cases.csv
../artifacts/protocol_rotation.csv

```

Columns in final_cases: ['case_id', 'case_position', 'block', 'loan_amnt', 'term', 'int_rate', 'log_annual_inc', 'dti', 'revol_util', 'home_ownership', 'purpose', 'credit_history_years', 'pred_prob', 'y_true', 'difficulty_tier', 'difficulty_score', 'correct', 'confidence', 'conf_bin', 'model_optimal']

Total cases: 18

Trials per participant: 18

Practice trials

```

[17]: # Practice trial 1: clear non-default (very low pred_prob, y_true=0)
practice_safe = df_pool[
    (df_pool['y_true'] == 0) &
    (df_pool['pred_prob'] < 0.06) &
    (df_pool['difficulty_tier'] == 'easy')
].sample(1, random_state=42)

# Practice trial 2: clear default (high pred_prob, y_true=1)
practice_risky = df_pool[
    (df_pool['y_true'] == 1) &
    (df_pool['pred_prob'] > 0.55) &
    (df_pool['difficulty_tier'] == 'hard')
].sample(1, random_state=42)

practice_cases = pd.concat([practice_safe, practice_risky]).
    ↪reset_index(drop=True)
practice_cases['case_position'] = [-2, -1] # negative = practice
practice_cases['block'] = 'practice'

print("=== PRACTICE TRIALS ===")
print(practice_cases[['case_id', 'pred_prob', 'y_true', 'difficulty_tier',

```



```

        'purpose', 'loan_amnt', 'int_rate']]].
    ↪to_string(index=False))

practice_cases.to_csv("../artifacts/practice_cases.csv", index=False)
print("\nSaved practice_cases.csv")

```

=== PRACTICE TRIALS ===

case_id	pred_prob	y_true	difficulty_tier	purpose	loan_amnt	int_rate
731479	0.058615	0	easy	moving	5000	8.24
252416	0.578106	1	hard	debt_consolidation	22000	19.89

Saved practice_cases.csv

Feature interpretation helpers

```

[18]: def dti_label(dti):
        if dti < 15: return "LOW"
        elif dti < 30: return "MODERATE"
        else: return "HIGH"

    def util_label(x):
        if x < 30: return "LOW"
        elif x < 70: return "MODERATE"
        else: return "HIGH"

    def income_label(x):
        if x < 30000: return "LOW"
        elif x < 100000: return "MEDIUM"
        else: return "HIGH"

```

Full case preview

INSTRUCTIONS

You are acting as a loan officer.

If you approve a loan and the borrower defaults → LARGE LOSS
If you reject a good loan → SMALL LOSS

Your goal is to make decisions that minimize total loss.

Your performance will be evaluated based on decision quality.

```

[19]: def render_case(row, protocol='ai_first', show_index=True):
        income = int(np.exp(row['log_annual_inc']))
        hist = round(row['credit_history_years'], 1)
        pred_pct = int(round(row['pred_prob'] * 100))
        tier = str(row['difficulty_tier']).upper()

```

```

# labels
dti_l = dti_label(row['dti'])
util_l = util_label(row['revol_util'])
inc_l = income_label(income)

header = ""
if show_index:
    header = (
        f"\n{'='*60}\n"
        f"CASE {int(row['case_position']):02d} of 18 | Block:␣
↪{row['block']} | [{tier}]\n"
        f"{'='*60}"
    )

profile = (
    f"\nLOAN APPLICATION\n"
    f"{'-'*44}\n"
    f"Loan amount      : ${row['loan_amnt']:>10,.0f}\n"
    f"Interest rate     : {row['int_rate']:>9.2f}%\n"
    f"Annual income      : ${income:>10,.0f} ({inc_l})\n"
    f"Debt-to-income     : {row['dti']:>9.1f} ({dti_l})\n"
    f"Revolving util.    : {row['revol_util']:>9.1f}% ({util_l})\n"
    f"Credit history     : {hist:>7.1f} years\n"
    f"Home ownership     : {str(row['home_ownership']):>12}\n"
    f"Loan purpose       : {str(row['purpose']):>12}\n"
    f"Term               : {str(row['term']):>12}\n"
    f"{'-'*44}"
)

# -----
# PROTOCOL LOGIC
# -----

if protocol == 'no_ai':
    decision_block = (
        f"\n\nDECISION TASK:\n"
        f"Would you APPROVE or REJECT this loan?\n"
        f"[Approve] [Reject]\n\n"
        f"(Optional) What probability do you assign to default?\n"
        f"[ 0% |-----| 100% ]\n\n"
        f"How confident are you? (0-100%)\n"
        f"[ 0% |-----| 100% ]\n"
    )

elif protocol == 'human_first':
    decision_block = (

```

```

        f"\n\nSTEP 1 - Initial decision:\n"
        f"Would you APPROVE or REJECT?\n"
        f"[Approve] [Reject]\n\n"
        f"(Optional) Initial probability:\n"
        f"[ 0% |-----| 100% ]\n"
        f"[SUBMIT]\n\n"
        f"--- AI revealed ---\n\n"
        f"AI ASSESSMENT:\n"
        f"Default probability: {pred_pct}%\n\n"
        f"STEP 2 - Final decision:\n"
        f"[Approve] [Reject]\n\n"
        f"(Optional) Final probability:\n"
        f"[ 0% |-----| 100% ]\n\n"
        f"Confidence (0-100%):\n"
        f"[ 0% |-----| 100% ]\n"
    )

else: # ai_first
    decision_block = (
        f"\n\nAI ASSESSMENT:\n"
        f"Default probability: {pred_pct}%\n\n"
        f"DECISION TASK:\n"
        f"Would you APPROVE or REJECT?\n"
        f"[Approve] [Reject]\n\n"
        f"(Optional) Probability estimate:\n"
        f"[ 0% |-----| 100% ]\n\n"
        f"Confidence (0-100%):\n"
        f"[ 0% |-----| 100% ]\n"
    )

    footer = f"\n{'-'*44}\n[SUBMIT]\n{'='*60}\n"

    print(header + profile + decision_block + footer)

```

Summary table

```

[20]: summary = final_cases[[
        'case_position', 'block', 'difficulty_tier',
        'correct', 'pred_prob', 'confidence',
        'y_true', 'model_optimal',
        'home_ownership', 'purpose', 'loan_amnt', 'int_rate'
    ]].copy()

summary.columns = [
    'pos', 'block', 'difficulty',
    'model_correct', 'ai_pred', 'ai_conf',
    'true_label', 'normative_align',

```

```

        'ownership', 'purpose', 'loan_amnt', 'int_rate'
    ]

    summary['ai_pred'] = summary['ai_pred'].round(3)
    summary['ai_conf'] = summary['ai_conf'].round(3)

    print("=== ALL 18 EXPERIMENTAL CASES ===\n")
    print(summary.to_string(index=False))

    print("\n=== KEY STATISTICS ===")
    print(f"Cases where model is correct : "
          f"{int(summary['model_correct'].sum())} / 18")
    print(f"Cases aligned with optimal : "
          f"{int(summary['normative_align'].sum())} / 18")
    print(f"True default rate in cases : "
          f"{summary['true_label'].mean():.3f}")
    print(f"Mean AI prediction : "
          f"{summary['ai_pred'].mean():.3f}")
    print(f"Mean AI confidence : "
          f"{summary['ai_conf'].mean():.3f}")
    print(f"Prediction range : "
          f"{summary['ai_pred'].min():.3f}, {summary['ai_pred'].max():.3f}")

    print("\n=== DESIGN BALANCE CHECK ===")
    print(pd.crosstab(
        summary['difficulty'],
        [summary['model_correct'], summary['normative_align']],
        rownames=['difficulty'],
        colnames=['model_correct', 'normative_align']
    ))

```

=== ALL 18 EXPERIMENTAL CASES ===

	pos	block	difficulty	model_correct	ai_pred	ai_conf	true_label
		normative_align	ownership		purpose	loan_amnt	int_rate
	1	block_1	easy	0	0.312	0.688	1
0		RENT	debt_consolidation	18000	10.91		
	2	block_1	easy	1	0.055	0.945	0
1		MORTGAGE	debt_consolidation	11000	5.32		
	3	block_1	hard	0	0.868	0.868	0
1		RENT	debt_consolidation	30225	30.75		
	4	block_1	hard	1	0.946	0.946	1
1		OWN	credit_card	18000	28.72		
	5	block_1	medium	0	0.234	0.766	1
0		RENT	moving	11450	15.02		
	6	block_1	medium	1	0.169	0.831	0
0		MORTGAGE	credit_card	4000	18.94		
	7	block_2	easy	0	0.333	0.667	1

0	MORTGAGE	small_business	40000	13.56		
	8 block_2	easy	1	0.075	0.925	0
1	MORTGAGE	debt_consolidation	10000	9.80		
	9 block_2	hard	0	0.843	0.843	0
1	RENT	debt_consolidation	25450	30.84		
	10 block_2	hard	1	0.923	0.923	1
1	RENT	debt_consolidation	13200	28.14		
	11 block_2	medium	0	0.267	0.733	1
0	RENT	debt_consolidation	8750	12.62		
	12 block_2	medium	1	0.189	0.811	0
0	MORTGAGE	credit_card	32000	13.99		
	13 block_3	easy	0	0.370	0.630	1
0	RENT	small_business	19200	15.04		
	14 block_3	easy	1	0.102	0.898	0
1	OWN	home_improvement	12000	7.99		
	15 block_3	hard	0	0.822	0.822	0
1	RENT	credit_card	35000	28.72		
	16 block_3	hard	1	0.901	0.901	1
1	RENT	debt_consolidation	14075	30.84		
	17 block_3	medium	0	0.289	0.711	1
0	OWN	debt_consolidation	8725	18.99		
	18 block_3	medium	1	0.210	0.790	0
0	MORTGAGE	debt_consolidation	2500	18.94		

=== KEY STATISTICS ===

Cases where model is correct : 9 / 18
 Cases aligned with optimal : 9 / 18
 True default rate in cases : 0.500
 Mean AI prediction : 0.439
 Mean AI confidence : 0.817
 Prediction range : [0.055, 0.946]

=== DESIGN BALANCE CHECK ===

model_correct	0	1		
normative_align	0	1	0	1
difficulty				
easy	3	0	0	3
hard	0	3	0	3
medium	3	0	3	0