

eda

March 29, 2026

0.0.1 Exploratory Data Analysis (EDA)

This analysis prepares a dataset for a **human–AI decision support system (DSS)** focused on loan approval decisions. The objective is not purely predictive performance, but the construction of a **controlled, interpretable decision environment** suitable for behavioral experimentation.

The EDA serves three primary purposes:

1. **Feature Space Construction**
 2. **Case Design for Experimentation**
 3. **Experimental Validity**
-

Key Design Principles

The entire pipeline is guided by the following principles:

- **Separation of model and interface**
 - Continuous features are used for modeling
 - Binned features are used for human interaction
 - **Reproducibility**
 - All transformations are deterministic and derived from training data only
 - **Interpretability over complexity**
 - Feature selection prioritizes clarity and usability over maximal predictive accuracy
 - **Alignment with human decision-making**
 - The dataset is structured to reflect how humans perceive and evaluate risk
-

Dataset Overview

We use a subset of Lending Club loan data filtered to a binary outcome:

- **Target variable:**
 $\text{target} = 1$ if loan is *Charged Off*, 0 if *Fully Paid*
- **Time span:**
Loans issued between 2007 and 2018
- **Default rate:**
Approximately 20%, indicating moderate class imbalance

This level of imbalance is preserved intentionally, as it reflects real-world decision conditions and is important for studying calibration and reliance behavior.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import spearmanr
from pandas.api.types import is_numeric_dtype
import os
import sys

np.random.seed(42)
pd.set_option('display.max_columns', None)
sns.set_style("whitegrid")

BASE_DIR = os.path.abspath(os.path.join(os.getcwd(), ".."))
sys.path.append(BASE_DIR)

from data_prep import load_and_clean, build_features

DATA_PATH = "../data/raw/loan.csv"

from feature_config import NUM_COLS, FEATURES, TARGET

print("Working directory:", os.getcwd())
```

Working directory: /Users/annaasatryan/Desktop/capstone/codes/notebooks

Data loading

```
[2]: df = load_and_clean(DATA_PATH)

print("\nShape:", df.shape)
print("Columns:", len(df.columns))
print("\nTime range:", df['issue_d'].min(), "→", df['issue_d'].max())
print("\nDefault rate:", df['target'].mean())

print("\nMissing (top 10):")
print(df.isna().mean().sort_values(ascending=False).head(10))
```

Shape: (1300076, 131)

Columns: 131

Time range: 2007-06-01 00:00:00 → 2018-12-01 00:00:00

Default rate: 0.2008374895006138

Missing (top 10):

id	1.000000
url	1.000000
member_id	1.000000
orig_projected_additional_accrued_interest	0.997359
payment_plan_start_date	0.995898
hardship_end_date	0.995898
hardship_type	0.995898
hardship_reason	0.995898
hardship_status	0.995898
deferral_term	0.995898
dtype: float64	

Domain-driven candidate features

```
[3]: df = build_features(df)

candidate_features = [
    # loan structure
    'loan_amnt',
    'term',
    'int_rate',
    'installment',

    # borrower profile
    # 'annual_inc',
    'log_annual_inc',
    'home_ownership',
    'purpose',

    # credit behavior
    'dti',
    'delinq_2yrs',
    'inq_last_6mths',
    'pub_rec_bankruptcies',
    'credit_history_years',
    'revol_util',
    'open_acc',
]

candidate_features = [c for c in candidate_features if c in df.columns]

print("\nInitial candidates:", candidate_features)
```

```
Initial candidates: ['loan_amnt', 'term', 'int_rate', 'installment',
'log_annual_inc', 'home_ownership', 'purpose', 'dti', 'delinq_2yrs',
'inq_last_6mths', 'pub_rec_bankruptcies', 'credit_history_years', 'revol_util',
'open_acc']
```

Missingness filter

```
[4]: missing = df[candidate_features].isna().mean().sort_values(ascending=False)
      print("\nMissingness:")
      print(missing)

      candidate_features = [c for c in candidate_features if missing[c] < 0.4]

      print("\nAfter missing filter:", candidate_features)
```

```
Missingness:
revol_util          6.184254e-04
pub_rec_bankruptcies 5.076626e-04
dti                 2.399860e-04
inq_last_6mths      7.691858e-07
loan_amnt           0.000000e+00
term                0.000000e+00
int_rate            0.000000e+00
installment         0.000000e+00
log_annual_inc      0.000000e+00
home_ownership      0.000000e+00
purpose             0.000000e+00
delinq_2yrs         0.000000e+00
credit_history_years 0.000000e+00
open_acc            0.000000e+00
dtype: float64
```

```
After missing filter: ['loan_amnt', 'term', 'int_rate', 'installment',
'log_annual_inc', 'home_ownership', 'purpose', 'dti', 'delinq_2yrs',
'inq_last_6mths', 'pub_rec_bankruptcies', 'credit_history_years', 'revol_util',
'open_acc']
```

Missingness across candidate features is negligible ($< 0.1\%$ for all retained variables). Simple imputation (median for numeric, mode for categorical variables) is sufficient and does not materially affect model calibration or downstream experimental validity.

Removing constant features

```
[5]: def is_constant(series):
      return series.nunique() <= 1

      candidate_features = [
          c for c in candidate_features
          if not is_constant(df[c])
      ]

      print("\nAfter constant filter:", candidate_features)
```

After constant filter: ['loan_amnt', 'term', 'int_rate', 'installment', 'log_annual_inc', 'home_ownership', 'purpose', 'dti', 'delinq_2yrs', 'inq_last_6mths', 'pub_rec_bankruptcies', 'credit_history_years', 'revol_util', 'open_acc']

A constant-feature filter is applied as a safeguard to remove degenerate variables with no variance.

Univariate analysis

```
[6]: def analyze_feature(df, feature, bins=5):
    df_temp = df[[feature, 'target']].dropna().copy()

    print(f"\n--- {feature} ---")

    if df_temp[feature].dtype in ['float64', 'int64']:
        df_temp['bin'] = pd.qcut(df_temp[feature], q=bins, duplicates='drop')

        agg = df_temp.groupby('bin', observed=True)['target'].mean()
        counts = df_temp['bin'].value_counts().sort_index()

        print(agg)

        x = range(len(agg))

        # Default rate plot
        plt.figure(figsize=(7,3))
        plt.plot(x, agg.values, marker='o')
        plt.xticks(x, agg.index.astype(str), rotation=45)
        plt.title(f"{feature} → default rate")
        plt.ylabel("Default rate")
        plt.tight_layout()
        plt.show()

        # Count plot
        plt.figure(figsize=(7,3))
        plt.bar(x, counts.values)
        plt.xticks(x, counts.index.astype(str), rotation=45)
        plt.title(f"{feature} → count")
        plt.tight_layout()
        plt.show()

    else:
        agg = df_temp.groupby(feature, observed=True)['target'].mean()
        counts = df_temp[feature].value_counts()

        print(agg)

        # Default rate
```

```

plt.figure(figsize=(7,3))
plt.bar(range(len(agg)), agg.values)
plt.xticks(range(len(agg)), agg.index.astype(str), rotation=45)
plt.title(f"{feature} → default rate")
plt.tight_layout()
plt.show()

# Count
plt.figure(figsize=(7,3))
plt.bar(range(len(counts)), counts.values)
plt.xticks(range(len(counts)), counts.index.astype(str), rotation=45)
plt.title(f"{feature} → count")
plt.tight_layout()
plt.show()

for f in candidate_features:
    analyze_feature(df, f)


```


--- loan_amnt ---

bin

(499.999, 6500.0]	0.159691
(6500.0, 10000.0]	0.176997
(10000.0, 15000.0]	0.208689
(15000.0, 21000.0]	0.228173
(21000.0, 40000.0]	0.234845

Name: target, dtype: float64






--- term ---


term

36 months 0.161015

60 months 0.325730

Name: target, dtype: float64







--- int_rate ---

bin

(5.308999999999999, 8.9]	0.067682
(8.9, 11.53]	0.138773
(11.53, 13.99]	0.194654
(13.99, 16.99]	0.251194
(16.99, 30.99]	0.360640

Name: target, dtype: float64







--- installment ---

bin

(4.928999999999999, 217.4]	0.154287
(217.4, 324.98]	0.191199
(324.98, 451.9]	0.217110
(451.9, 640.58]	0.219550
(640.58, 1719.83]	0.222087

Name: target, dtype: float64







--- log_annual_inc ---

bin

(-0.001, 10.645]	0.236732
(10.645, 10.951]	0.217210
(10.951, 11.212]	0.204306
(11.212, 11.513]	0.185133
(11.513, 16.213]	0.158039

Name: target, dtype: float64







```

--- home_ownership ---
home_ownership
ANY          0.194757
MORTGAGE     0.173030
NONE         0.145833
OTHER        0.194245
OWN          0.207565
RENT         0.233765
Name: target, dtype: float64

```







--- purpose ---

```

purpose
car                0.146388
credit_card        0.170267
debt_consolidation 0.212567
home_improvement   0.178531
house              0.217167
major_purchase     0.187235
medical            0.219064
moving             0.234409
other              0.211706
small_business     0.297468
vacation           0.192396
Name: target, dtype: float64

```







--- dti ---

bin

(-1.001, 10.5]	0.149530
(10.5, 15.33]	0.168552
(15.33, 19.98]	0.192321
(19.98, 25.68]	0.220839
(25.68, 999.0]	0.273090

Name: target, dtype: float64







```

--- delinq_2yrs ---
bin
(-0.001, 39.0]    0.200837
Name: target, dtype: float64

```







```

--- inq_last_6mths ---
bin
(-0.001, 1.0]    0.191730
(1.0, 8.0]      0.250167
Name: target, dtype: float64

```







```

--- pub_rec_bankruptcies ---
bin
(-0.001, 12.0]    0.200853
Name: target, dtype: float64

```







--- credit_history_years ---

bin

(2.9970000000000003, 10.42]	0.222339
(10.42, 13.328]	0.213994
(13.328, 16.498]	0.197919
(16.498, 21.582]	0.190151
(21.582, 83.25]	0.179733

Name: target, dtype: float64





--- revol_util ---

bin

(-0.001, 29.2] 0.158371


(29.2, 45.1] 0.191146


(45.1, 59.4] 0.208644

(59.4, 75.1] 0.218329

(75.1, 892.3] 0.227809

Name: target, dtype: float64





--- open_acc ---

bin

(-0.001, 7.0] 0.186035


(7.0, 9.0] 0.193504


(9.0, 12.0] 0.202872

(12.0, 15.0] 0.206650

(15.0, 90.0] 0.217280

Name: target, dtype: float64





Categorical features vs target

```
[7]: print("\n=== Categorical Features vs Target ===")

cat_cols = ['purpose', 'home_ownership', 'term']

for col in cat_cols:
    print(f"\n--- {col} ---")
    summary = (
        df.groupby(col, observed=True)['target']
        .agg(['mean', 'count'])
        .sort_values('mean', ascending=False)
    )
    print(summary)
```

=== Categorical Features vs Target ===

--- purpose ---

	mean	count
purpose		
small_business	0.297468	15010
moving	0.234409	9172
medical	0.219064	15023
house	0.217167	6967
debt_consolidation	0.212567	757591
other	0.211706	74934
vacation	0.192396	8732
major_purchase	0.187235	28328
home_improvement	0.178531	84495
credit_card	0.170267	285704
car	0.146388	14120

```

--- home_ownership ---
              mean  count
home_ownership
RENT          0.233765  515761
OWN           0.207565  139561
ANY           0.194757    267
OTHER         0.194245    139
MORTGAGE      0.173030  644300
NONE          0.145833    48

```

```

--- term ---
              mean  count
term
60 months    0.325730  314316
36 months    0.161015  985760

```

Correlation check


```

[8]: numeric_features = [
      c for c in candidate_features
      if is_numeric_dtype(df[c])
    ]

corr = df[numeric_features].corr()

plt.figure(figsize=(14,4))
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title("Correlation Matrix")
plt.show()

```



The correlation matrix is used to identify redundancy between features rather than predictive

strength. A strong correlation is observed between loan_amnt and installment (0.95), confirming that installment is functionally derived from loan structure variables. No other pairs exhibit high correlation, indicating low multicollinearity among retained features.

Spearman

```
[9]: results = []

for col in NUM_COLS:
    valid = df[[col, 'target']].dropna()
    corr, pval = spearmanr(valid[col], valid['target'])
    results.append((col, corr, pval))

spearman_df = pd.DataFrame(results,
    columns=['feature', 'spearman_corr', 'p_value'])

print(spearman_df.sort_values('spearman_corr', key=abs, ascending=False))
```

	feature	spearman_corr	p_value
1	int_rate	0.254142	0.0
3	dti	0.108542	0.0
0	loan_amnt	0.070226	0.0
2	log_annual_inc	-0.067729	0.0
5	revol_util	0.060362	0.0
4	credit_history_years	-0.039248	0.0

Spearman rank correlation is used to assess monotonic relationships between features and default risk. The strongest relationship is observed for int_rate (0.25), followed by dti and loan_amnt, indicating moderate predictive signal.

These results are used jointly with univariate analysis to guide feature selection.

Feature decisions

```
[10]: redundant_features = ['installment']
low_signal_features = ['delinq_2yrs', 'pub_rec_bankruptcies']
parsimony_exclusions = ['inq_last_6mths', 'open_acc']

remove_features = set(
    redundant_features +
    low_signal_features +
    parsimony_exclusions
)

final_features = [
    f for f in candidate_features
    if f not in remove_features
]

final_features
```

```
[10]: ['loan_amnt',  
      'term',  
      'int_rate',  
      'log_annual_inc',  
      'home_ownership',  
      'purpose',  
      'dti',  
      'credit_history_years',  
      'revol_util']
```

Feature Selection and Justification

Each feature is retained only if it satisfies all of the following:

1. **Available at decision time**
Features must be observable before the lending decision (no post-outcome leakage).
2. **Interpretability**
Features must be understandable to human participants in the DSS.
3. **Stable relationship with target**
The feature should show a consistent (ideally monotonic) relationship with default risk in EDA.
4. **Non-redundancy**
Features that are highly correlated or functionally derived from others are removed.

1. Redundant Feature

- **installment removed**

Reason: - Highly correlated with `loan_amnt` - Functionally determined by (`loan_amnt`, `term`, `int_rate`)

Implication: - Adds no independent information - Keeping it would introduce multicollinearity and distort interpretability

2. Near-No-Signal Features

- **delinq_2yrs removed**
- **pub_rec_bankruptcies removed**

Reason: - Flat or near-constant default rate across bins in EDA - No meaningful discriminative power

Implication: - Including them adds noise without improving predictive or experimental value

3. Moderate-Signal but Excluded for Parsimony The following features showed some predictive signal but were excluded to reduce feature space complexity and cognitive burden in the decision support system:

- **inq_last_6mths removed**
 - Limited resolution (few distinct values)
 - Weak and unstable relationship with default risk
- **open_acc removed**
 - Weak monotonic trend
 - Overlaps conceptually with broader credit capacity indicators

Reason: - These features provide marginal incremental value relative to stronger predictors - Including them increases dimensionality without meaningful gain in interpretability or experimental clarity

4. Final Retained Features The final feature set includes only variables with: - clear directional signal - intuitive meaning - non-redundant information content

Loan structure - loan_amnt, term, int_rate

Borrower profile - log_annual_inc, home_ownership, purpose

Credit behavior - dti, credit_history_years, revol_util

We deliberately favor a **compact, high-signal feature set** to:

- ensure participants can process information without overload
- preserve interpretability of decisions
- isolate the effect of interaction protocols rather than feature complexity

This aligns with the objective of evaluating **human-AI interaction**, not maximizing predictive performance.

Temporal split

```
[11]: train = df[df['issue_d'] < "2016-01-01"].copy()
      test  = df[df['issue_d'] >= "2016-01-01"].copy()

print("\nTrain default:", train['target'].mean())
print("Test default:", test['target'].mean())
print("Shift:", test['target'].mean() - train['target'].mean())
```

Train default: 0.18448195087671354

Test default: 0.228741694836228

Shift: 0.04425974395951446

Temporal distribution shift introduces distribution shift that may affect calibration

Binning on train

```
[12]: bin_edges = {}

_, edges = pd.qcut(
    train['log_annual_inc'],
    q=5,
    retbins=True,
    duplicates='drop'
)
bin_edges['income'] = edges

_, edges = pd.qcut(
    train['loan_amnt'],
    q=4,
    retbins=True,
    duplicates='drop'
)
bin_edges['loan'] = edges

_, edges = pd.qcut(
    train['int_rate'],
    q=4,
    retbins=True,
    duplicates='drop'
)
bin_edges['rate'] = edges

_, edges = pd.qcut(
    train['dti'],
    q=4,
    retbins=True,
    duplicates='drop'
)
bin_edges['dti'] = edges

[13]: def apply_bins(df, bin_edges):
    df = df.copy()

    # Extend edges to capture out-of-range values
    def open_edges(edges):
        e = edges.copy()
        e[0] = -np.inf
        e[-1] = np.inf
        return e

    df['income_bin'] = pd.cut(df['log_annual_inc'],
```

```

        bins=open_edges(bin_edges['income']),
        include_lowest=True)
df['loan_bin'] = pd.cut(df['loan_amnt'],
        bins=open_edges(bin_edges['loan']),
        include_lowest=True)
df['rate_bin'] = pd.cut(df['int_rate'],
        bins=open_edges(bin_edges['rate']),
        include_lowest=True)
df['dti_bin'] = pd.cut(df['dti'],
        bins=open_edges(bin_edges['dti']),
        include_lowest=True)

return df

# Apply everywhere
train = apply_bins(train, bin_edges)
test = apply_bins(test, bin_edges)
df = apply_bins(df, bin_edges)

```

Bin reliability

```

[14]: for col in ['income_bin', 'loan_bin', 'rate_bin', 'dti_bin']:
        print(f"\n=== {col} ===")
        print(
            train
            .groupby(col, observed=True)['target']
            .agg(['mean', 'count'])
        )

```

=== income_bin ===

	mean	count
income_bin		
(-inf, 10.645]	0.215146	174305
(10.645, 10.915]	0.202768	154842
(10.915, 11.184]	0.190546	167057
(11.184, 11.493]	0.169284	161557
(11.493, inf]	0.142886	161891

=== loan_bin ===

	mean	count
loan_bin		
(-inf, 8000.0]	0.145539	222380
(8000.0, 12000.0]	0.177131	189137
(12000.0, 20000.0]	0.204499	233409
(20000.0, inf]	0.215263	174726

=== rate_bin ===

	mean	count
--	------	-------

```
rate_bin
(-inf, 9.76]      0.068912  205698
(9.76, 12.99]    0.142919  220397
(12.99, 15.88]   0.209188  190412
(15.88, inf]     0.323439  203145
```

```
=== dti_bin ===
```

```
          mean  count
dti_bin
(-inf, 11.78]  0.137226  204997
(11.78, 17.5]  0.163951  205147
(17.5, 23.76]  0.193332  204777
(23.76, inf]   0.243517  204729
```

Stability check

```
[15]: for col in ['income_bin', 'loan_bin', 'rate_bin', 'dti_bin']:
      print(f"\n=== {col} ===")

      print("Train:")
      print(
          train
          .groupby(col, observed=True)['target']
          .mean()
      )

      print("Test:")
      print(
          test
          .groupby(col, observed=True)['target']
          .mean()
      )
```

```
=== income_bin ===
```

```
Train:
```

```
income_bin
(-inf, 10.645]    0.215146
(10.645, 10.915]  0.202768
(10.915, 11.184]  0.190546
(11.184, 11.493]  0.169284
(11.493, inf]     0.142886
```

```
Name: target, dtype: float64
```

```
Test:
```

```
income_bin
(-inf, 10.645]    0.278072
(10.645, 10.915]  0.246001
(10.915, 11.184]  0.230566
(11.184, 11.493]  0.214468
```

```
(11.493, inf]      0.185864
Name: target, dtype: float64
```

```
=== loan_bin ===
```

```
Train:
```

```
loan_bin
```

```
(-inf, 8000.0]      0.145539
(8000.0, 12000.0]   0.177131
(12000.0, 20000.0]  0.204499
(20000.0, inf]      0.215263
```

```
Name: target, dtype: float64
```

```
Test:
```

```
loan_bin
```

```
(-inf, 8000.0]      0.187304
(8000.0, 12000.0]   0.221538
(12000.0, 20000.0]  0.250832
(20000.0, inf]      0.266383
```

```
Name: target, dtype: float64
```

```
=== rate_bin ===
```

```
Train:
```

```
rate_bin
```

```
(-inf, 9.76]        0.068912
(9.76, 12.99]       0.142919
(12.99, 15.88]      0.209188
(15.88, inf]        0.323439
```

```
Name: target, dtype: float64
```

```
Test:
```

```
rate_bin
```

```
(-inf, 9.76]        0.094704
(9.76, 12.99]       0.188895
(12.99, 15.88]      0.265489
(15.88, inf]        0.373259
```

```
Name: target, dtype: float64
```

```
=== dti_bin ===
```

```
Train:
```

```
dti_bin
```

```
(-inf, 11.78]       0.137226
(11.78, 17.5]       0.163951
(17.5, 23.76]       0.193332
(23.76, inf]        0.243517
```

```
Name: target, dtype: float64
```

```
Test:
```

```
dti_bin
```

```
(-inf, 11.78]       0.176635
(11.78, 17.5]       0.200471
(17.5, 23.76]       0.234060
```

```
(23.76, inf]      0.295487
Name: target, dtype: float64
```

Case difficulty for experiment

```
[16]: # -----
# Design choices:
# - Model-independent (uses empirical base rates)
# - Multi-dimensional (rate + dti + income)
# - Smoothed (reduces noise in sparse bins)
# - Train-only estimation (no leakage)
# -----

BIN_COLS = ['rate_bin', 'dti_bin', 'income_bin']

# 1. Compute bin statistics on TRAIN ONLY
bin_stats = (
    train
    .groupby(BIN_COLS, observed=True)
    .agg(
        bin_count=(TARGET, 'size'),
        bin_mean=(TARGET, 'mean')
    )
)

# 2. Apply Bayesian smoothing (shrinkage)
global_mean = train[TARGET].mean()
alpha = 50 #

bin_stats['smoothed_rate'] = (
    (bin_stats['bin_count'] * bin_stats['bin_mean'] + alpha * global_mean)
    / (bin_stats['bin_count'] + alpha)
)

bin_rates = bin_stats['smoothed_rate']

# 3. Mapping function (robust to unseen bins)
def map_base_rate(df_source, bin_rates, global_mean):
    idx = pd.MultiIndex.from_frame(df_source[BIN_COLS])
    mapped = idx.map(bin_rates)
    return mapped.fillna(global_mean) # fallback for unseen bins

# 4. Map to TRAIN and FULL DATA
train['case_base_rate'] = map_base_rate(train, bin_rates, global_mean)
df['case_base_rate'] = map_base_rate(df, bin_rates, global_mean)

# 5. Difficulty score (distance from 0.5)
```

```

# Max uncertainty at 0.5 → hardest
df['difficulty_score'] = 1 - np.abs(df['case_base_rate'] - 0.5) / 0.5

# Numerical safety
df['difficulty_score'] = df['difficulty_score'].clip(0, 1)

# 6. Stable tiering (fixed thresholds)
df['difficulty_tier'] = pd.cut(
    df['difficulty_score'],
    bins=[0, 0.33, 0.66, 1],
    labels=['easy', 'medium', 'hard'],
    include_lowest=True
)

# 7. Validation checks (MANDATORY)
print("\n=== Difficulty sanity check ===")
print(df.groupby('difficulty_tier', observed=True)[TARGET].mean())

print("\n=== Counts per tier ===")
print(df['difficulty_tier'].value_counts())

print("\n=== Bin coverage ===")
print("Missing rate:", df['case_base_rate'].isna().mean())

```

```

=== Difficulty sanity check ===
difficulty_tier
easy      0.115213
medium    0.253935
hard      0.384850
Name: target, dtype: float64

```

```

=== Counts per tier ===
difficulty_tier
easy      640042
medium    509118
hard      150916
Name: count, dtype: int64

```

```

=== Bin coverage ===
Missing rate: 0.0

```

Multiple formulations of case difficulty were evaluated, including distance from empirical base rate and distribution-based (quantile) approaches. The final definition was selected based on two criteria: (1) monotonic separation of observed default rates across difficulty tiers, and (2) sufficient sample size per tier for stable estimation.

The hard tier is smaller in size due to the rarity of high-uncertainty cases in the data; this is preserved to maintain realism rather than enforcing artificial balance.

DSS structure

```
[17]: model_df = df[FEATURES + [TARGET]].dropna().copy()

experiment_df = df[
    ['loan_bin', 'rate_bin', 'income_bin', 'dti_bin',
     'home_ownership', 'purpose',
     TARGET, 'difficulty_tier']
].dropna().copy()

print("\nModel dataset:", model_df.shape)
print("Experiment dataset:", experiment_df.shape)
```

Model dataset: (1298960, 10)

Experiment dataset: (1299764, 8)

Final dataset export

```
[18]: df_final = df[FEATURES + [TARGET, 'issue_d']].copy()

df_final['case_base_rate'] = df['case_base_rate']
df_final['difficulty_score'] = df['difficulty_score']
df_final['difficulty_tier'] = df['difficulty_tier']

df_final = df_final.dropna().reset_index(drop=True)
df_final['case_id'] = df_final.index

df_final.to_csv("../data/processed/loan_v1.csv", index=False)
```

```
[19]: df_final.head()
```

```
[19]:   loan_amnt    term  int_rate  dti  revol_util  home_ownership  \
0    30000  36 months   22.35  30.46        37.0        MORTGAGE
1    40000  60 months   16.14  50.53        64.5        MORTGAGE
2    20000  36 months    7.56  18.92        29.9        MORTGAGE
3     4500  36 months   11.31   4.64        15.3           RENT
4     8425  36 months   27.27  12.37        65.7        MORTGAGE

      purpose  log_annual_inc  credit_history_years  target  \
0  debt_consolidation    11.512935           6.915811      0
1      credit_card    10.714440           9.500342      0
2      credit_card    11.512935          19.830253      0
3      credit_card    10.558439          15.000684      0
4      credit_card    13.017005          21.166324      0

      issue_d  case_base_rate  difficulty_score  difficulty_tier  case_id
0  2018-12-01      0.331939         0.663877             hard         0
1  2018-12-01      0.390871         0.781741             hard         1
```

2	2018-12-01	0.065519	0.131038	easy	2
3	2018-12-01	0.143847	0.287694	easy	3
4	2018-12-01	0.264996	0.529992	medium	4