

modeling

April 17, 2026

0.0.1 Modeling notebook

LendingClub Default Prediction for DSS

Objective: Calibrated probability model for human-AI experiment

```
[1]: import os
import sys
import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import HistGradientBoostingClassifier
from sklearn.calibration import CalibratedClassifierCV, calibration_curve
from sklearn.metrics import (
    roc_auc_score, brier_score_loss, log_loss,
    confusion_matrix, ConfusionMatrixDisplay
)
from sklearn.base import clone
from scipy.stats import ttest_rel
from xgboost import XGBClassifier

BASE_DIR = os.path.abspath(os.path.join(os.getcwd(), ".."))
sys.path.append(BASE_DIR)

from feature_config import FEATURES, NUM_COLS, CAT_COLS, TARGET

np.random.seed(42)
pd.set_option('display.max_columns', None)
sns.set_style("whitegrid")
```

Data loading

```
[2]: DATA_PATH = "../../../data/processed/loan_v1.csv"
df = pd.read_csv(DATA_PATH)
df['issue_d'] = pd.to_datetime(df['issue_d'], errors='coerce')

print("Shape:", df.shape)
print("Default rate:", round(df[TARGET].mean(), 4))
print("Date range:", df['issue_d'].min(), "→", df['issue_d'].max())
print("\nFeatures used:", FEATURES)
```

Shape: (1298960, 15)

Default rate: 0.2008

Date range: 2007-06-01 00:00:00 → 2018-12-01 00:00:00

Features used: ['loan_amnt', 'term', 'int_rate', 'dti', 'revol_util', 'home_ownership', 'purpose', 'log_annual_inc', 'credit_history_years']

Temporal splits

Split strategy:

- train → chronological 70% of pre-2016 (model fitting)
- cal → first half of remaining 30% (calibration only)
- val → second half of remaining 30% (model selection)
- test → post-2016 (final reporting, used once)

Calibration is isolated from validation to prevent isotonic regression from memorizing its fitting set (ECE → 0 artifact). Model selection is done on val. Test is never used for selection.

```
[3]: train_full = df[df['issue_d'] < "2016-01-01"].copy()
test      = df[df['issue_d'] >= "2016-01-01"].copy()

cutoff_cal = train_full['issue_d'].quantile(0.70)
train      = train_full[train_full['issue_d'] < cutoff_cal].copy()
cal_val    = train_full[train_full['issue_d'] >= cutoff_cal].copy()

cutoff_val = cal_val['issue_d'].quantile(0.50)
cal        = cal_val[cal_val['issue_d'] < cutoff_val].copy()
val        = cal_val[cal_val['issue_d'] >= cutoff_val].copy()

y_train = train[TARGET]
y_cal   = cal[TARGET]
y_val   = val[TARGET]
y_test  = test[TARGET]

# Features ONLY for modeling
```

```

X_train = train[FEATURES].copy()
X_cal   = cal[FEATURES].copy()
X_val   = val[FEATURES].copy()
X_test  = test[FEATURES].copy()

# Keep case_id separately (for export only)
case_id_train = train['case_id'].copy()
case_id_cal   = cal['case_id'].copy()
case_id_val   = val['case_id'].copy()
case_id_test  = test['case_id'].copy()

print(f"Train : {train.shape[0]:>8,} | default rate: {y_train.mean():.4f}")
print(f"Cal   : {cal.shape[0]:>8,} | default rate: {y_cal.mean():.4f}")
print(f"Val   : {val.shape[0]:>8,} | default rate: {y_val.mean():.4f}")
print(f"Test  : {test.shape[0]:>8,} | default rate: {y_test.mean():.4f}")

```

```

Train : 554,109 | default rate: 0.1765
Cal   : 126,135 | default rate: 0.2031
Val   : 138,972 | default rate: 0.1995
Test  : 479,744 | default rate: 0.2288

```

Outlier capping (train-derived)

DTI has extreme outliers. Clip at 99.9th percentile of train. Same threshold applied to all splits — no leakage.

```

[4]: dti_upper = X_train['dti'].quantile(0.999)
for X in [X_train, X_cal, X_val, X_test]:
    X['dti'] = X['dti'].clip(upper=dti_upper)

print(f"DTI clipped at {dti_upper:.2f} (99.9th percentile of train)")

```

DTI clipped at 39.53 (99.9th percentile of train)

Preprocessing pipeline

```

[5]: preprocessor = ColumnTransformer(
    transformers=[
        ('num', Pipeline([
            ('imputer', SimpleImputer(strategy='median')),
            ('scaler', StandardScaler())
        ]), NUM_COLS),
        ('cat', Pipeline([
            ('imputer', SimpleImputer(strategy='most_frequent')),
            ('onehot', OneHotEncoder(
                drop='first',
                handle_unknown='ignore',
                sparse_output=False
            ))
        ]), CAT_COLS)
    )

```

```
]
)
```

Metric function

ECE uses fixed equal-width bins. 10 bins chosen as standard bias–variance tradeoff for calibration estimation. Fixed bins ensure ECE is comparable across subgroups, time periods, and models. Quantile bins shift with each distribution making cross-subset comparison invalid.

```
[6]: def compute_ece(y_true, y_prob, n_bins=10):
    y_true = np.array(y_true)
    y_prob = np.array(y_prob)
    bins = np.linspace(0, 1, n_bins + 1)
    ece = 0.0
    for i in range(n_bins):
        if i < n_bins - 1:
            mask = (y_prob >= bins[i]) & (y_prob < bins[i + 1])
        else:
            mask = (y_prob >= bins[i]) & (y_prob <= bins[i + 1])
        if mask.sum() == 0:
            continue
        acc = y_true[mask].mean()
        conf = y_prob[mask].mean()
        ece += (mask.sum() / len(y_true)) * abs(acc - conf)
    return ece

def evaluate(y_true, y_prob, label=""):
    return {
        'Model' : label,
        'AUC' : round(roc_auc_score(y_true, y_prob), 4),
        'Brier' : round(brier_score_loss(y_true, y_prob), 4),
        'LogLoss': round(log_loss(y_true, y_prob), 4),
        'ECE' : round(compute_ece(y_true, y_prob), 4),
    }

# Bootstrap uses fixed seed for reproducibility
def bootstrap_ci(y_true, y_prob, metric_fn, n=1000, alpha=0.05, seed=42):
    y_true = np.array(y_true)
    y_prob = np.array(y_prob)
    rng = np.random.default_rng(seed)
    N = len(y_true)
    vals = []
    for _ in range(n):
        idx = rng.choice(N, N, replace=True)
        vals.append(metric_fn(y_true[idx], y_prob[idx]))
    lo = round(np.percentile(vals, 100 * alpha / 2), 4)
    hi = round(np.percentile(vals, 100 * (1 - alpha / 2)), 4)
```

```
return lo, hi
```

Model definitions

- Three models: interpretable baseline (Logistic) + twogradient boosting variants (XGB, HGB).
- Hyperparameters are fixed (not grid-searched)
- the goal is a defensible, calibrated DSS model, not maximizing AUC.

```
[7]: models = {
    'Logistic': Pipeline([
        ('preprocessor', preprocessor),
        ('model', LogisticRegression(
            C=0.1,
            max_iter=3000,
            random_state=42
        ))
    ]),
    'XGB': Pipeline([
        ('preprocessor', preprocessor),
        ('model', XGBClassifier(
            n_estimators=300,
            max_depth=4,
            learning_rate=0.05,
            subsample=0.8,
            colsample_bytree=0.8,
            eval_metric='logloss',
            random_state=42,
            verbosity=0
        ))
    ]),
    'HGB': Pipeline([
        ('preprocessor', preprocessor),
        ('model', HistGradientBoostingClassifier(
            max_iter=300,
            max_depth=4,
            learning_rate=0.05,
            random_state=42
        ))
    ])
}
```

Resampled training stability analysis

```
[8]: N_RUNS = 10
rng = np.random.default_rng(42)

records = []

for run in range(N_RUNS):
```

```

print(f"\n=== BOOTSTRAP RUN {run+1}/{N_RUNS} ===")

# -----
# 1. Bootstrap sample (TRAIN ONLY)
# -----
idx = rng.choice(len(X_train), len(X_train), replace=True)
X_train_boot = X_train.iloc[idx]
y_train_boot = y_train.iloc[idx]

for name, base_pipe in models.items():

    # ensure fresh model
    pipe = clone(base_pipe)

    # -----
    # 2. Fit on bootstrapped data
    # -----
    pipe.fit(X_train_boot, y_train_boot)

    # -----
    # 3. Calibration
    # -----
    iso = CalibratedClassifierCV(pipe, method='isotonic', cv='prefit')
    sig = CalibratedClassifierCV(pipe, method='sigmoid', cv='prefit')

    iso.fit(X_cal, y_cal)
    sig.fit(X_cal, y_cal)

    # -----
    # 4. Select calibration (Brier)
    # -----
    probs_val_iso = iso.predict_proba(X_val)[: , 1]
    probs_val_sig = sig.predict_proba(X_val)[: , 1]

    brier_iso = brier_score_loss(y_val, probs_val_iso)
    brier_sig = brier_score_loss(y_val, probs_val_sig)

    if brier_iso <= brier_sig:
        chosen = iso
    else:
        chosen = sig

    # -----
    # 5. Evaluate on TEST
    # -----
    probs_test = chosen.predict_proba(X_test)[: , 1]

```

```

auc = roc_auc_score(y_test, probs_test)
brier = brier_score_loss(y_test, probs_test)
ece = compute_ece(y_test, probs_test)

records.append({
    'run': run,
    'model': name,
    'AUC': auc,
    'Brier': brier,
    'ECE': ece
})

robust_df = pd.DataFrame(records)

print("\n=== ROBUSTNESS SUMMARY ===\n")

summary = robust_df.groupby('model')[['AUC', 'Brier', 'ECE']].agg(['mean', 'std'])
print(summary)

```

```

=== BOOTSTRAP RUN 1/10 ===

=== BOOTSTRAP RUN 2/10 ===

=== BOOTSTRAP RUN 3/10 ===

=== BOOTSTRAP RUN 4/10 ===

=== BOOTSTRAP RUN 5/10 ===

=== BOOTSTRAP RUN 6/10 ===

=== BOOTSTRAP RUN 7/10 ===

=== BOOTSTRAP RUN 8/10 ===

=== BOOTSTRAP RUN 9/10 ===

=== BOOTSTRAP RUN 10/10 ===

=== ROBUSTNESS SUMMARY ===

```

	AUC		Brier		ECE	
model	mean	std	mean	std	mean	std
HGB	0.696435	0.000436	0.162783	0.000076	0.032156	0.000459
Logistic	0.691300	0.000318	0.163946	0.000078	0.033068	0.000201
XGB	0.696626	0.000422	0.162891	0.000104	0.033801	0.000741

Statistical comparison

```
[9]: def paired_test(df, m1, m2, metric):
    x = df[df['model']==m1][metric].values
    y = df[df['model']==m2][metric].values

    stat, p = ttest_rel(x, y)
    diff = x - y

    print(f"\n{m1} vs {m2} ({metric})")
    print(f"Mean difference: {diff.mean():.6f}")
    print(f"Std difference : {diff.std():.6f}")
    print(f"t-stat: {stat:.4f}, p-value: {p:.6f}")

    # Compare models
    paired_test(robust_df, 'Logistic', 'HGB', 'Brier')
    paired_test(robust_df, 'Logistic', 'XGB', 'Brier')
    paired_test(robust_df, 'HGB', 'XGB', 'Brier')
```

```
Logistic vs HGB (Brier)
Mean difference: 0.001163
Std difference : 0.000060
t-stat: 58.4699, p-value: 0.000000
```

```
Logistic vs XGB (Brier)
Mean difference: 0.001055
Std difference : 0.000112
t-stat: 28.3278, p-value: 0.000000
```

```
HGB vs XGB (Brier)
Mean difference: -0.000108
Std difference : 0.000072
t-stat: -4.4598, p-value: 0.001578
```

Model comparison across logistic regression, XGBoost, and histogram gradient boosting showed minimal differences in predictive performance. Bootstrap analysis confirmed high stability across training resamples, with Brier score differences between models on the order of 0.001. While these differences are statistically significant due to large sample size, they are practically negligible.

Train + calibrate all models

Pipeline: 1. Fit base model on `X_train` 2. Fit isotonic and sigmoid calibration on `X_cal` (`cv='prefit'`) 3. Select calibration method using VAL BRIER (proper scoring rule) 4. Store val Brier for model selection (NOT test) 5. Store test probabilities for final reporting only

```
[10]: X_train_model = X_train.copy()
      X_cal_model   = X_cal.copy()
      X_val_model    = X_val.copy()
```



```

X_test_model = X_test.copy()

trained = {}

for name, base_pipe in models.items():
    print(f"\nTraining {name}...")

    # ensure fresh model
    pipe = clone(base_pipe)

    # 1. Fit base model
    pipe.fit(X_train_model, y_train)

    # 2. Calibration models
    iso = CalibratedClassifierCV(pipe, method='isotonic', cv='prefit')
    sig = CalibratedClassifierCV(pipe, method='sigmoid', cv='prefit')

    iso.fit(X_cal_model, y_cal)
    sig.fit(X_cal_model, y_cal)

    # 3. Predictions on validation
    probs_val_iso = iso.predict_proba(X_val_model)[: , 1]
    probs_val_sig = sig.predict_proba(X_val_model)[: , 1]

    # 4. Select calibration via Brier
    brier_iso_val = brier_score_loss(y_val, probs_val_iso)
    brier_sig_val = brier_score_loss(y_val, probs_val_sig)

    if abs(brier_iso_val - brier_sig_val) < 1e-6:
        chosen = sig
        method = 'sigmoid'
        probs_val = probs_val_sig
    elif brier_iso_val < brier_sig_val:
        chosen = iso
        method = 'isotonic'
        probs_val = probs_val_iso
    else:
        chosen = sig
        method = 'sigmoid'
        probs_val = probs_val_sig

    # diagnostics
    ece_iso_val = compute_ece(y_val, probs_val_iso)
    ece_sig_val = compute_ece(y_val, probs_val_sig)

    # 5. Test predictions
    probs_test = chosen.predict_proba(X_test_model)[: , 1]

```

```

val_brier = brier_score_loss(y_val, probs_val)

trained[name] = {
    'model': chosen,
    'probs_val': probs_val,
    'probs_test': probs_test,
    'val_brier': round(val_brier, 4),
    'cal_method': method
}

print(f"  Calibration : {method}")
print(f"  Val Brier    : iso={brier_iso_val:.4f} sig={brier_sig_val:.4f}")
print(f"  Val ECE      : iso={ece_iso_val:.4f} sig={ece_sig_val:.4f}")

print("\nAll models trained and calibrated.")

```

Training Logistic...

```

Calibration : isotonic
Val Brier    : iso=0.1419 sig=0.1420
Val ECE      : iso=0.0080 sig=0.0123

```

Training XGB...

```

Calibration : isotonic
Val Brier    : iso=0.1413 sig=0.1423
Val ECE      : iso=0.0108 sig=0.0282

```

Training HGB...

```

Calibration : sigmoid
Val Brier    : iso=0.1412 sig=0.1412
Val ECE      : iso=0.0091 sig=0.0091

```

All models trained and calibrated.

Model selection

We DO compute validation Brier for all models, but final model is FIXED to Logistic based on:
 - negligible performance differences (~ 0.001 Brier) - robustness analysis - need for interpretability and experimental clarity

```

[11]: val_briers = {name: d['val_brier'] for name, d in trained.items()}

print("=== VALIDATION BRIER (model comparison) ===\n")
for name, b in val_briers.items():
    print(f"  {name}: {b:.4f}")

# FIXED FINAL MODEL

```

```

FINAL_MODEL_NAME = 'Logistic'
final_model      = trained['Logistic']['model']
final_probs      = trained['Logistic']['probs_test']

print(f"\nSelected: {FINAL_MODEL_NAME} (fixed for interpretability &
↪equivalence)")

# Prepare output for DSS / case selection
test_output = X_test.copy() # includes case_id
test_output['y_true'] = y_test.values
test_output['pred_prob'] = final_probs

```

=== VALIDATION BRIER (model comparison) ===

```

Logistic: 0.1419
XGB: 0.1413
HGB: 0.1412

```

Selected: Logistic (fixed for interpretability & equivalence)

Feature importance

```

[12]: # Get the fitted pipeline
final_pipe = trained['Logistic']['model']

# The calibrated model wraps the base pipeline
base_pipe = final_pipe.calibrated_classifiers_[0].estimator
preprocessor_fitted = base_pipe.named_steps['preprocessor']
lr_model = base_pipe.named_steps['model']

# Get feature names after preprocessing
num_names = NUM_COLS
cat_encoder = preprocessor_fitted.named_transformers_['cat'].
↪named_steps['onehot']
cat_names = list(cat_encoder.get_feature_names_out(CAT_COLS))
all_names = num_names + cat_names

# Coefficients
coef_df = pd.DataFrame({
    'feature': all_names,
    'coefficient': lr_model.coef_[0]
}).sort_values('coefficient', key=abs, ascending=False)

print("=== LOGISTIC REGRESSION COEFFICIENTS ===\n")
print(coef_df.to_string(index=False))

# Save for paper
coef_df.to_csv("../artifacts/lr_coefficients.csv", index=False)

```


```

# Plot top 10
fig, ax = plt.subplots(figsize=(8, 5))
top = coef_df.head(10)
colors = ['tomato' if c > 0 else 'steelblue' for c in top['coefficient']]
ax.barh(range(len(top)), top['coefficient'].values, color=colors)
ax.set_yticks(range(len(top)))
ax.set_yticklabels(top['feature'].values)
ax.set_xlabel("Coefficient (log-odds)")
ax.set_title("Top 10 Features by Absolute Coefficient")
ax.invert_yaxis()
plt.tight_layout()
plt.savefig("../visualizations/feature_importance.png", dpi=150,
            bbox_inches='tight')
plt.show()

```

=== LOGISTIC REGRESSION COEFFICIENTS ===

	feature	coefficient
	term_60 months	0.604922
	home_ownership_MORTGAGE	-0.531648
	purpose_small_business	0.527611
	int_rate	0.441987
	home_ownership_OWN	-0.419191
	home_ownership_RENT	-0.278563
	log_annual_inc	-0.189749
	purpose_medical	0.184769
	purpose_home_improvement	0.163031
	purpose_moving	0.157153
	dti	0.141294
	purpose_debt_consolidation	0.121528
	loan_amnt	0.074156
	home_ownership_NONE	-0.073928
	purpose_other	0.055939
	purpose_credit_card	0.053597
	purpose_major_purchase	0.049495
	purpose_house	-0.048972
	purpose_vacation	0.023197
	home_ownership_OTHER	-0.016785
	credit_history_years	-0.014121
	revol_util	0.004081



Model comparison table

```
[13]: results = []

for name, d in trained.items():
    m = evaluate(y_test, d['probs_test'], label=name)

    m['Calibration'] = d['cal_method']
    m['Val_Brier'] = d['val_brier']

    # Bootstrap confidence intervals
    auc_lo, auc_hi = bootstrap_ci(
        y_test.values,
        d['probs_test'],
        roc_auc_score
    )

    brier_lo, brier_hi = bootstrap_ci(
        y_test.values,
        d['probs_test'],
        brier_score_loss
    )

    m['AUC_95CI'] = f"[{auc_lo}, {auc_hi}]"
    m['Brier_95CI'] = f"[{brier_lo}, {brier_hi}]"
```

```

        results.append(m)

results_df = pd.DataFrame(results).set_index('Model')

print("\n=== MODEL COMPARISON - TEST SET (FINAL REPORTING) ===\n")
print(
    results_df[
        ['AUC', 'Brier', 'LogLoss', 'ECE',
         'Calibration', 'Val_Brier',
         'AUC_95CI', 'Brier_95CI']
    ].to_string()
)

# Effect sizes
delta_hgb = results_df.loc['Logistic', 'Brier'] - results_df.loc['HGB', 'Brier']
delta_xgb = results_df.loc['Logistic', 'Brier'] - results_df.loc['XGB', 'Brier']

print("\n=== EFFECT SIZE (Brier differences) ===")
print(f"Logistic - HGB: {delta_hgb:.6f}")
print(f"Logistic - XGB: {delta_xgb:.6f}")

print(f"\nFinal model: {FINAL_MODEL_NAME}")
print("Test metrics:")
print(evaluate(y_test, final_probs, label=FINAL_MODEL_NAME))

```

=== MODEL COMPARISON - TEST SET (FINAL REPORTING) ===

	AUC	Brier	LogLoss	ECE	Calibration	Val_Brier	
AUC_95CI	Brier_95CI						
Model							
Logistic	0.6914	0.1639	0.5132	0.0328	isotonic	0.1419	[0.6897, 0.6932]
XGB	0.6972	0.1628	0.4998	0.0340	isotonic	0.1413	[0.6955, 0.699]
HGB	0.6974	0.1626	0.4978	0.0326	sigmoid	0.1412	[0.6956, 0.6991]

=== EFFECT SIZE (Brier differences) ===

Logistic - HGB: 0.001300

Logistic - XGB: 0.001100

Final model: Logistic

Test metrics:

```
{
  'Model': 'Logistic',
  'AUC': np.float64(0.6914),
  'Brier': np.float64(0.1639),
  'LogLoss': 0.5132,
  'ECE': np.float64(0.0328)
}
```


Logistic regression was selected for its interpretability and transparency. Using a simple and interpretable model minimizes confounding factors introduced by model complexity, ensuring that observed behavioral effects can be attributed to interaction structure rather than differences in model behavior.

Calibration curves

```
[14]: fig, axes = plt.subplots(1, 3, figsize=(15, 4))

for ax, (name, d) in zip(axes, trained.items()):
    prob_true, prob_pred = calibration_curve(
        y_test, d['probs_test'], n_bins=10
    )
    ece = compute_ece(y_test.values, d['probs_test'])
    ax.plot(prob_pred, prob_true, marker='o', label='Model')
    ax.plot([0, 1], [0, 1], '--', color='gray', label='Perfect')
    ax.set_title(f"{name} - ECE {ece:.4f}")
    ax.set_xlabel("Mean predicted probability")
    ax.set_ylabel("Fraction of positives")
    ax.legend(fontsize=9)

plt.suptitle("Calibration Curves - Test Set", y=1.02)
plt.tight_layout()
plt.savefig("../visualizations/calibration_curves.png",
            dpi=150, bbox_inches='tight')
plt.show()
```



ECE by difficulty tier

```
[15]: test_with_diff = test[['case_id', 'difficulty_tier']].copy()
test_with_diff['pred_prob'] = final_probs
test_with_diff['y_true'] = y_test.values

print("=== ECE BY DIFFICULTY TIER ===\n")
for tier in ['easy', 'medium', 'hard']:
    mask = test_with_diff['difficulty_tier'] == tier
```

```

sub_y = test_with_diff.loc[mask, 'y_true'].values
sub_p = test_with_diff.loc[mask, 'pred_prob'].values
ece = compute_ece(sub_y, sub_p)
brier = brier_score_loss(sub_y, sub_p)
print(f"{tier:8s}: ECE={ece:.4f}, Brier={brier:.4f}, n={mask.sum():,}")

```

=== ECE BY DIFFICULTY TIER ===

```

easy      : ECE=0.0248, Brier=0.1155, n=241,081
medium    : ECE=0.0399, Brier=0.2032, n=175,439
hard      : ECE=0.0496, Brier=0.2391, n=63,224

```

Probability distribution

```


[16]: probs_arr = np.array(final_probs)

plt.figure(figsize=(8, 3))
plt.hist(probs_arr, bins=60, edgecolor='none', alpha=0.8)
plt.axvline(0.5, color='red', linestyle='--', label='0.5 threshold')
plt.axvline(0.3, color='orange', linestyle='--', label='0.3 threshold')
plt.title(f"Predicted Probability Distribution - {FINAL_MODEL_NAME}")
plt.xlabel("Predicted P(default)")
plt.ylabel("Count")
plt.legend()
plt.tight_layout()
plt.savefig("../visualizations/prob_distribution.png",
            dpi=150, bbox_inches='tight')
plt.show()

print("=== PROBABILITY DISTRIBUTION DIAGNOSTICS ===\n")
print(f"Sharpness (variance)      : {np.var(probs_arr):.4f}")
print(f"Mean predicted P(default): {probs_arr.mean():.4f}")
print(f"Actual default rate       : {y_test.mean():.4f}")
print(f"Calibration gap (mean)    : "
      f"{abs(probs_arr.mean() - y_test.mean()):.4f}")
print(f"\nFraction predicted > 0.5 : {(probs_arr > 0.5).mean():.4f}")
print(f"Fraction predicted > 0.3 : {(probs_arr > 0.3).mean():.4f}")
print(f"Fraction predicted > 0.2 : {(probs_arr > 0.2).mean():.4f}")

print(f"\nPrediction range distribution:")
bins_diag = [0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 1.0]
labels_diag = ['0-0.1', '0.1-0.2', '0.2-0.3',
               '0.3-0.4', '0.4-0.5', '0.5-0.6', '0.6+']
counts = (pd.cut(probs_arr, bins=bins_diag, labels=labels_diag)
          .value_counts()
          .sort_index())
for label, count in counts.items():
    print(f" {label}: {count:>8,} ({count/len(probs_arr)*100:.1f}%)"

```

=== PROBABILITY DISTRIBUTION DIAGNOSTICS ===

Sharpness (variance) : 0.0224
 Mean predicted P(default): 0.2151
 Actual default rate : 0.2288
 Calibration gap (mean) : 0.0137

Fraction predicted > 0.5 : 0.0783
 Fraction predicted > 0.3 : 0.2307
 Fraction predicted > 0.2 : 0.4377

Prediction range distribution:

0-0.1:	111,402	(23.2%)
0.1-0.2:	158,317	(33.0%)
0.2-0.3:	99,331	(20.7%)
0.3-0.4:	49,041	(10.2%)
0.4-0.5:	24,095	(5.0%)
0.5-0.6:	24,975	(5.2%)
0.6+:	12,565	(2.6%)

Interpretation: Probability Compression

The model produces low-variance probabilities concentrated in [0.1–0.3]. This results in a low-signal decision environment, where cases are not trivially separable. Such conditions are desirable for studying human–AI interaction, as they increase reliance on probabilistic reasoning and amplify the role of interaction structure in decision-making.

Confusion matrices at multiple thresholds

The DSS presents probabilities to participants — it does NOT make binary approve/reject decisions. Threshold-based analysis is included here solely as a diagnostic to:

- (a) document class imbalance behavior
- (b) show recall-precision tradeoff at different operating points


Primary evaluation metrics are Brier and ECE (probability-based).

```
[17]: thresholds = [0.3, 0.4, 0.5]
fig, axes = plt.subplots(1, 3, figsize=(15, 4))

for ax, thresh in zip(axes, thresholds):
    preds = (final_probs >= thresh).astype(int)
    cm = confusion_matrix(y_test, preds)
    disp = ConfusionMatrixDisplay(cm, display_labels=['Paid', 'Default'])
    disp.plot(ax=ax, colorbar=False, cmap='Blues')

    tn, fp, fn, tp = cm.ravel()
    recall = tp / (tp + fn)
    precision = tp / (tp + fp) if (tp + fp) > 0 else 0
    ax.set_title(
        f"Threshold = {thresh}\n"
        f"Recall: {recall:.3f} Precision: {precision:.3f}"
    )

plt.suptitle(
    f"Confusion Matrices - {FINAL_MODEL_NAME} - Test Set\n"
    "Diagnostic only. DSS uses predicted probabilities, "
    "not binary decisions.",
    y=1.02
)
plt.tight_layout()
plt.savefig("../visualizations/confusion_matrices.png",
            dpi=150, bbox_inches='tight')
plt.show()
```



Feature distribution shift

```
[18]: shift_features = ['int_rate', 'dti', 'log_annual_inc']
fig, axes = plt.subplots(1, 3, figsize=(15, 4))
```


```

for ax, col in zip(axes, shift_features):
    ax.hist(train[col].dropna(), bins=50, alpha=0.6, density=True,
            label=f'Train  $\mu$ ={train[col].mean():.2f}')
    ax.hist(test[col].dropna(), bins=50, alpha=0.6, density=True,
            label=f'Test  $\mu$ ={test[col].mean():.2f}')
    ax.set_title(f"{col} - Distribution Shift")
    ax.set_xlabel(col)
    ax.set_ylabel("Density")
    ax.legend(fontsize=8)

plt.suptitle("Feature Distribution: Train vs Test (post-2016)", y=1.02)
plt.tight_layout()
plt.savefig("../visualizations/feature_shift.png",
            dpi=150, bbox_inches='tight')
plt.show()

print("=== FEATURE DISTRIBUTION SHIFT ===\n")
for col in shift_features:
    delta = test[col].mean() - train[col].mean()
    print(f"{col:20s} train  $\mu$ ={train[col].mean():.3f} "
          f"test  $\mu$ ={test[col].mean():.3f}  $\Delta$ ={delta:+.3f}")

```



=== FEATURE DISTRIBUTION SHIFT ===

int_rate	train	=13.548	test	=13.467	Δ =-0.081
dti	train	=17.524	test	=18.766	Δ =+1.242
log_annual_inc	train	=11.058	test	=11.115	Δ =+0.056

A moderate distribution shift is observed in debt-to-income ratio (Δ +1.24), indicating mild covariate shift between training and test periods. This further motivates the use of calibrated probabilities rather than threshold-based decisions.

Export test predictions (final, with features + difficulty)

```
[19]: test_output = X_test.copy()
```

```

# attach identifiers
test_output['case_id'] = case_id_test.values

# attach labels
test_output['y_true'] = y_test.values
test_output['pred_prob'] = final_probs

test_output = test_output.merge(
    test[['case_id', 'difficulty_score', 'difficulty_tier']],
    on='case_id',
    how='left'
)

# reorder
FEATURES_ORDER = list(X_test.columns)

test_output = test_output[[
    'case_id',
    *FEATURES_ORDER,
    'y_true',
    'pred_prob',
    'difficulty_score',
    'difficulty_tier'
]]

# checks
assert not test_output['difficulty_tier'].isna().any(), "Missing difficulty!"

# save
test_output.to_parquet("../artifacts/test_predictions.parquet", index=False)

print("Saved FULL test_predictions.parquet")

```

Saved FULL test_predictions.parquet